

RAID types

RAID (Redundant Array of Inexpensive Disks or Drives, or Redundant Array of Independent Disks).

Purposes:

performance improvement (max. write read speed),
increased data security against hard drive failure.


Hardware RAID = RAID controller

or


Software RAID

RAID types


RAID 0




RAID 1



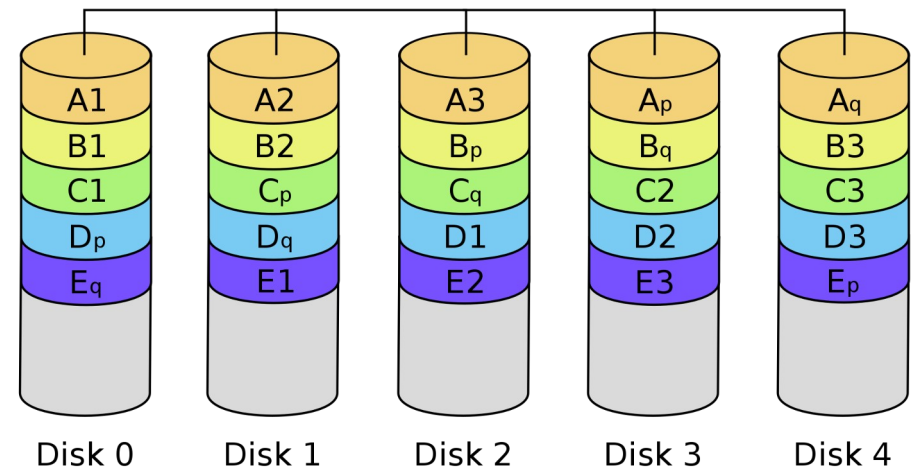
RAID 2 (Hamming code)



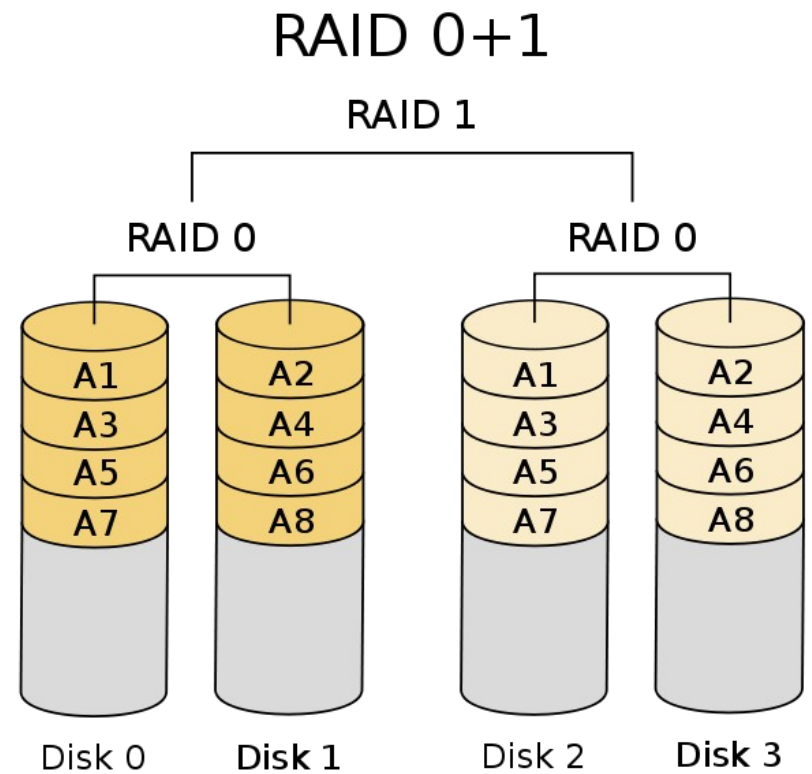
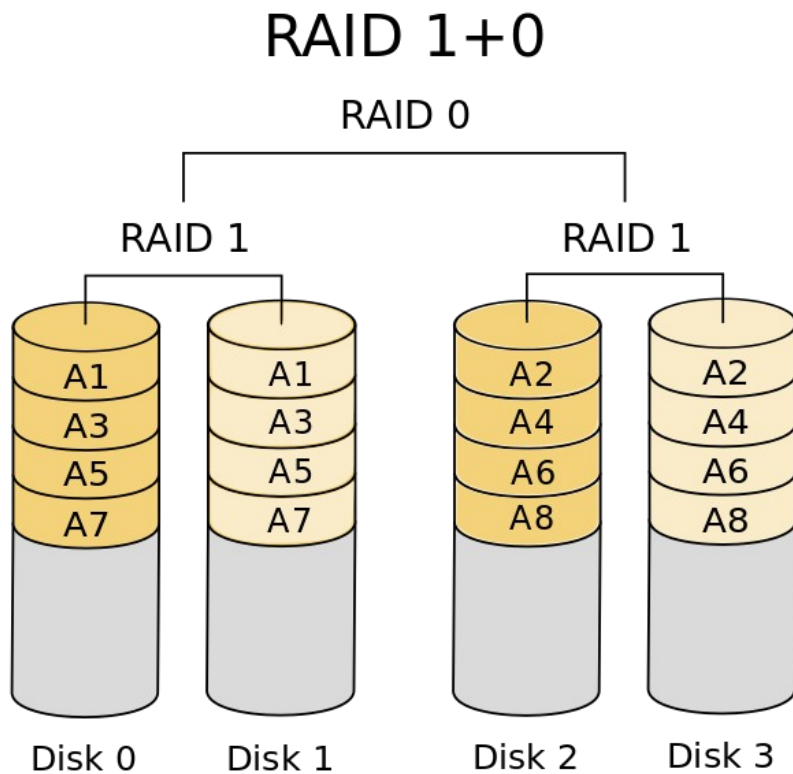
RAID 5



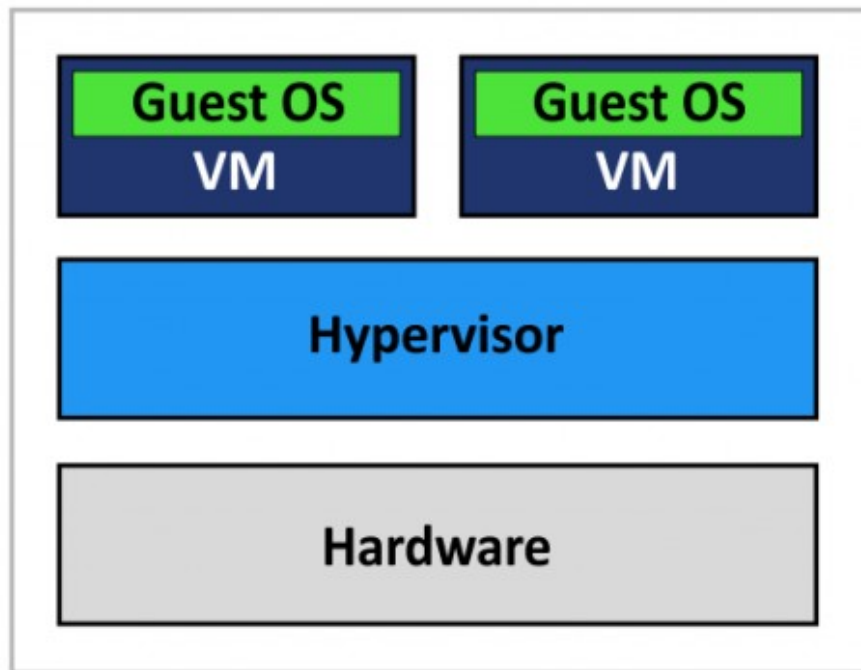
RAID 6



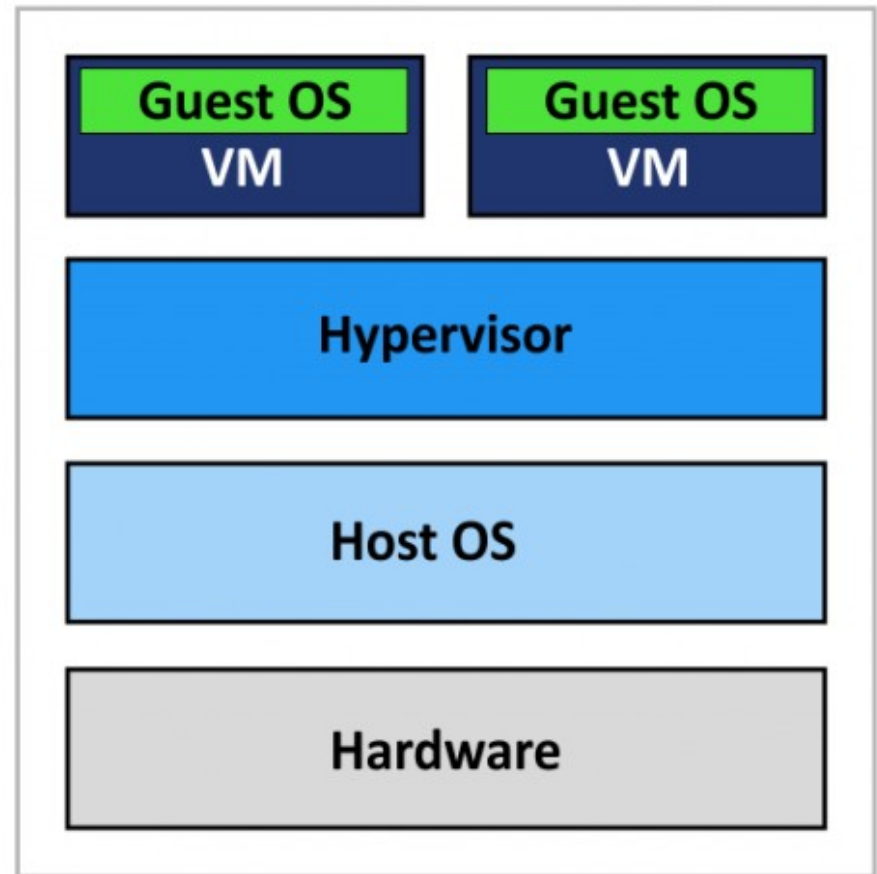
Mixed RAID types principle



Virtual machine approach

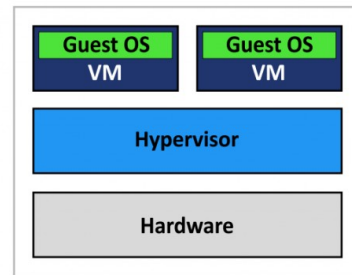


**Type 1 Hypervisor
(Bare-Metal Architecture)**

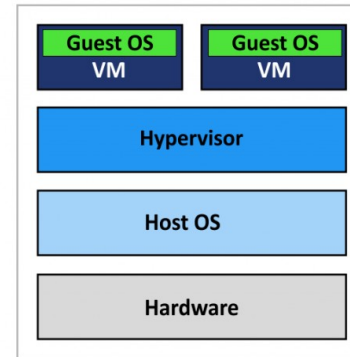


**Type 2 Hypervisor
(Hosted Architecture)**

Virtual machine approach



Type 1 Hypervisor
(Bare-Metal Architecture)

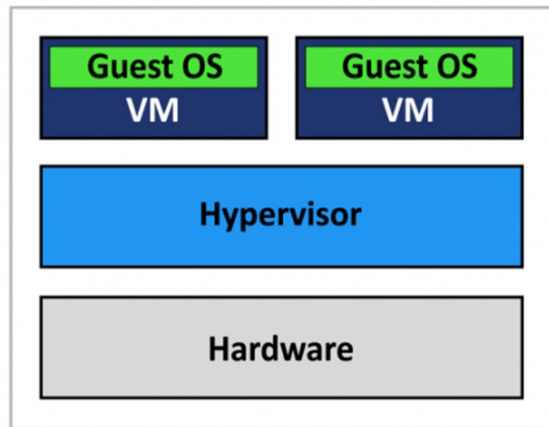


Type 2 Hypervisor
(Hosted Architecture)

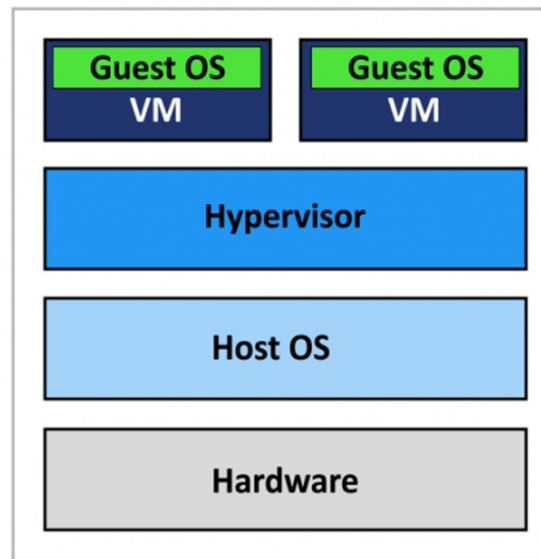
Vmware ESXi
Microsoft Hyper-V Server
Citrix XenServer
KVM

VMware Workstation
Virtualbox

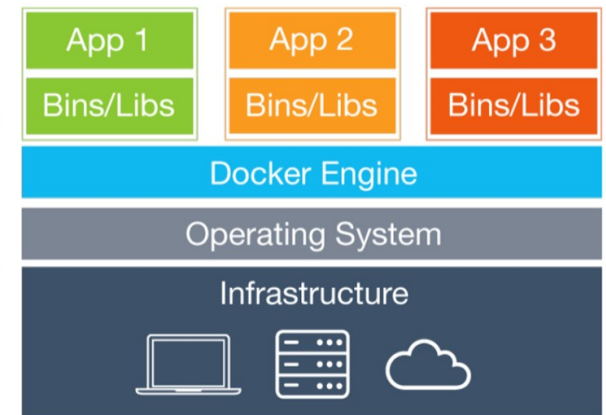
VM vs Container



Type 1 Hypervisor
(Bare-Metal Architecture)

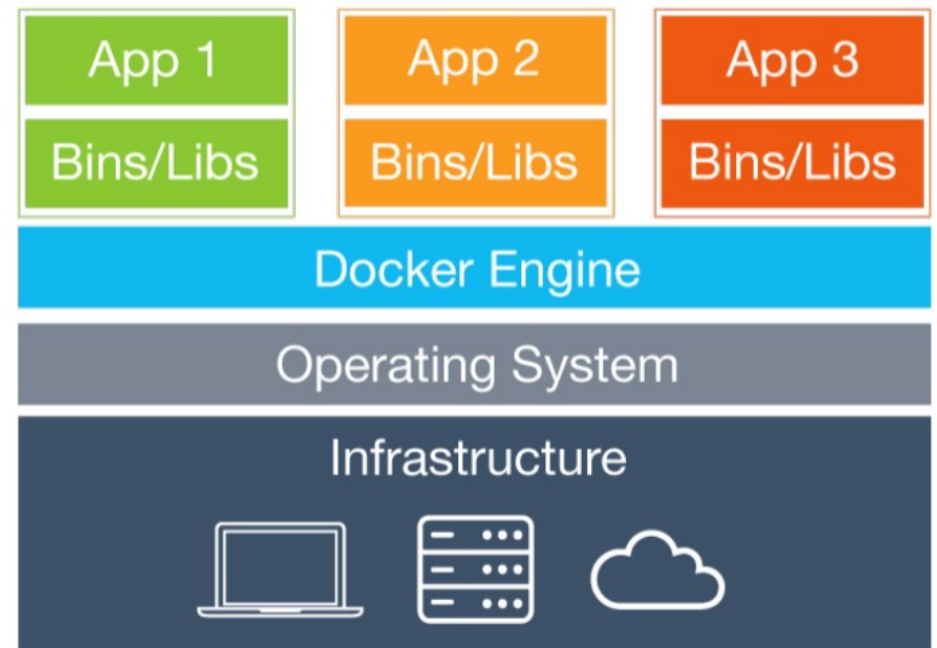
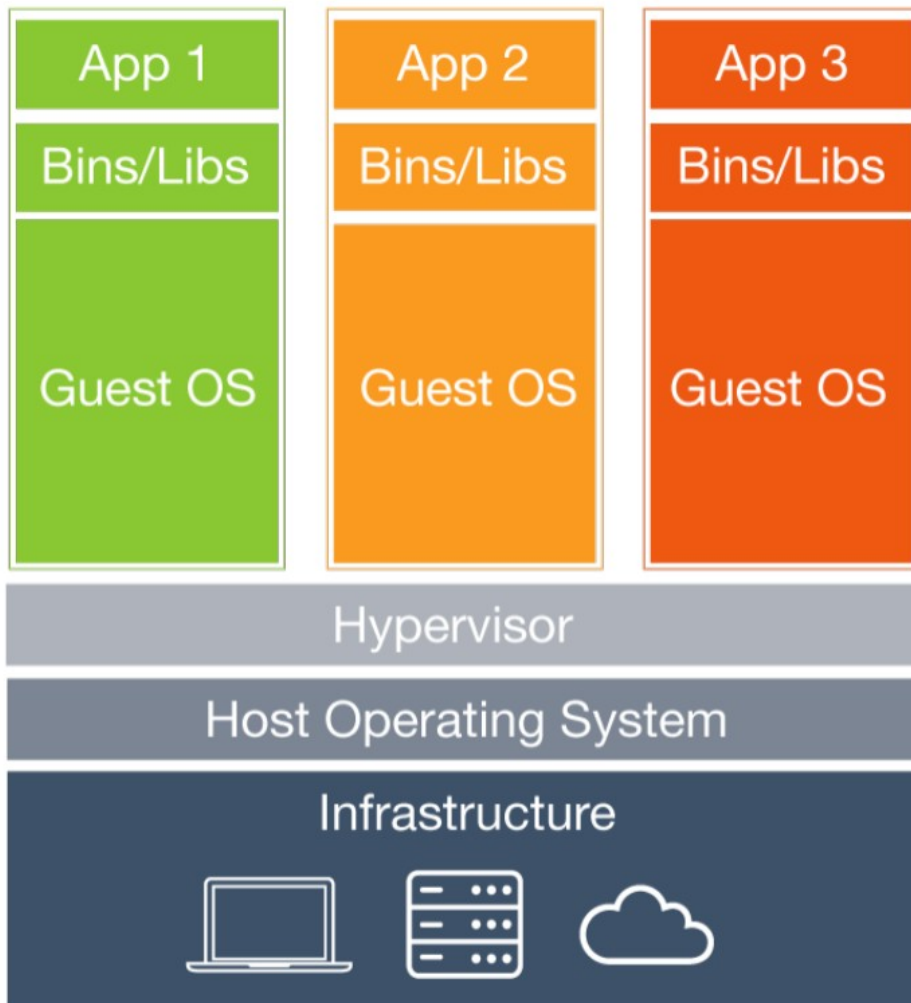


Type 2 Hypervisor
(Hosted Architecture)

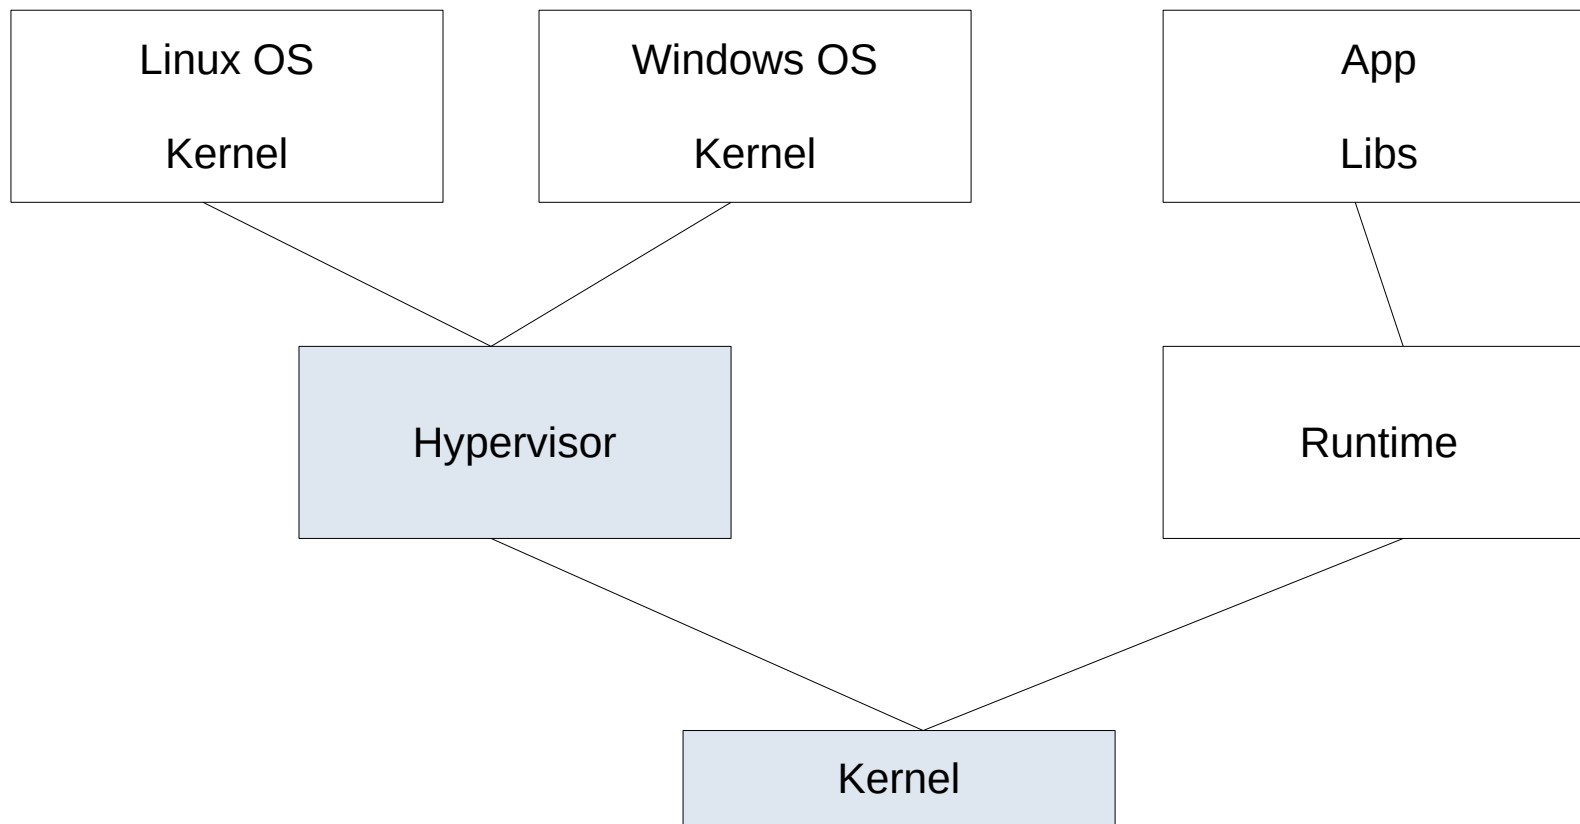


Docker

VM vs Container



VM vs Container



Containers

Linux containers are ordinary processes, which isolated by Linux kernel tools:

- * Namespaces for resource isolation
- * Control Groups for resource constraints
- * Permissions, Capabilities, SELinux, AppArmor and other Linux security constraints

All Linux processes are started as containers.

- * `cat /proc/PID/cgroup` will see the Cgroups the process is in
- * `cat /proc/self/attr/current` will show SELinux labels
- * `ls /proc/PID/ns` will show a list of namespaces the process is in

Container different is that it is started from a container image. So, container runtime is the operating system part that manages the cgroups, SELinux labels and namespaces and starts the container from an image

Container is

- * A container is a running instance of an image
- * The image contains the application code, language runtime and libraries
- * External libraries such as libc are typically provided by the host operating system, but in a container is included in the images
- * The container image is a read-only instance of the application that just needs to be started
- * While starting a container, it adds a writable layer on the top to store any changes that are made while working with the container

What is an Image?

- * A container image is a TAR file that combines the following
 - * The container root file system: a directory that looks like the standard root of the operating system, but presented as a mount namespace
 - * Metadata: a JSON file that specifies how to run that root file system, including all settings required to get to a functional container (entrypoint, environment variables and more)
- * Container images are layered: you can install additional content, add a new JSON file and store the differences in a new TAR file
- * Images are standardized by the OCI

- * Container images are typically shared through public registries, or by sharing mechanisms to build them easily, such as Dockerfile
- * The Docker container image format has become the de facto standard image format
- * Open Container Initiative (OCI) has standardized the Docker container image format

Container Image Layers

- * Docker images are made up of a series of filesystem layers
- * Each layer adds, removes or modifies files from the preceding layer in the filesystem
- * This is called an overlay filesystem, and different overlay filesystems exist, like aufs, overlay and overlay2
- * By using these different image layers, and pointing to other image layers in a smart way, it's easy to build container images that have support for multiple versions of vital components
- * Apart from the different layers, container images have a container configuration file that provides instructions on how to run the container


App	
Rails3	Rails4
Ruby	
Alpine	

Current Linux Namespaces

- * **pid**: isolation for processes
- * **user**: presents isolated users only existing inside the namespace
- * **mnt**: the modern alternative to the chroot jail
- * **net**: provides isolated networking
- * **ipc**: inter-process communication, used for sharing memory
- * **uts**: UNIX time sharing, which allows processes to have their own hostname and domain name
- * **cgroup**: providing isolated root directories for each cgroup
- * **time**: used to virtualize the clock of a system

Container Engines. Container Runtimes

- * In the early days of containers (2014-2015), containers were started and managed by a container engine
- * A container engine offers multiple components
 - * A core component that runs the containers
 - * Optionally a daemon that controls the containers
 - * A command line interface
- * Common container engines include and included
 - * Docker
 - * LXC
 - * Podman
 - * System-nspawn



- * While container engines developed further, container runtimes were split off as separate projects
- * By splitting of container runtimes into different open source projects, it became the easier to focus on developing and standardizing them
- * Containerd became the common runtime in Docker environments
- * CRI-o is the common runtime in Podman environments

Docker Essentials

<https://docs.docker.com/get-started/>